VentureFlow Intelligence: A Multimodal Deep Learning Architecture for Startup Success Prediction

Annanahmed Shaikh
Wentworth Institute of Technology
Boston, Massachusetts, USA
shaikha4@wit.edu

Abstract

Early-stage venture capital investment decisions are dominated by subjective judgment, with an estimated 75–80% of funded startups failing to return their investors' capital. At the Seed and Series A stages, the absence of revenue history, customer traction, or product-market fit data renders traditional quantitative models ineffective. This thesis presents VentureFlow Intelligence, a decision-support platform that combines semantic analysis of startup business narratives with structured financial signals to predict startup success. Our system employs a Dual-Encoder architecture: Branch A uses the Sentence-BERT model `all-MiniLM-L6-v2` to compute cosine similarity between a startup's text description and a learned "Unicorn Centroid," while Branch B applies log-normalized funding data as a financial stability check. Evaluated on a longitudinal dataset of 47,833 global startups, the Dual-Encoder Fusion model achieves an ROC-AUC of 0.74, representing a 10% absolute improvement over the best tabular baseline (Random Forest at 0.675 AUC). An 11-step weight ablation study confirms the optimal 60/40 financial-to-semantic weighting. Out-of-distribution testing on 2024–2025 generative AI companies (xAI, Anthropic) demonstrates robust temporal generalization. The platform is deployed as a production Streamlit dashboard with a Star Schema data warehouse and real-time news integration.

CCS Concepts

• **Computing methodologies → Machine learning**.

Keywords

Venture Capital, Startup Prediction, Sentence Embeddings, Multimodal Fusion, NLP, Decision Support Systems

ACM Reference Format:
Annanahmed Shaikh. 2026. VentureFlow Intelligence: A Multimodal Deep Learning Architecture for Startup Success Prediction. In *Proceedings of Wentworth Institute of Technology (WIT MSDS '26).* ACM, New York, NY, USA, 12 pages. https://doi.org/10.1145/nnnnnnn.nnnnnnn

1 Introduction

Every year, billions of dollars flow into early-stage startups, yet the vast majority of these investments fail. Venture capital analysts face a fundamental information problem: the companies they must evaluate at the earliest stages have no revenue, no customers, and often no finished product. This thesis addresses that challenge by building a system that reads what a startup *says about itself*—its business description—and combines that textual signal with basic financial data to predict whether the company is likely to succeed. The resulting platform, called VentureFlow Intelligence, achieved meaningfully stronger predictive performance than traditional machine-learning approaches and is deployed as an interactive dashboard for institutional use. This section introduces the economic context, the information gap we target, and the research objectives that guided our work across two capstone semesters.

1.1 The Power Law in Venture Capital

The venture capital asset class follows a power-law return distribution rather than the Gaussian distribution observed in public equities [14]. A small fraction of portfolio companies generates the vast majority of returns; the rest return little or nothing. Empirical industry data indicates that between 75% and 80% of all venture-backed startups fail to return their investors' initial capital [6, 11]. The entire asset class survives economically because a handful of "unicorn" outcomes—companies reaching billion-dollar valuations—compensate for the losses on everything else.

This asymmetry creates a decision environment where false negatives (missing a future unicorn) carry far greater cost than false positives (funding a company that fails). Standard classification approaches that optimize for overall accuracy are therefore poorly suited to venture capital, because they tend to suppress the rare positive class in favor of the dominant negative class [16]. Instead, predictive models for venture investment must be evaluated primarily on their ability to surface high-potential outliers—a task well-suited to anomaly-detection-style scoring methods.

1.2 The Pre-Traction Information Vacuum

At the Seed and Series A stages, startups typically lack the financial track record that traditional valuation models require [18]. Features that are central to public-equity analysis—such as trailing twelve-month (TTM) revenue, customer acquisition cost (CAC), price-to-earnings (P/E) ratio, and debt-to-equity ratio—are simply unavailable for a pre-revenue company. In formal terms, the standard feature vector for a mature enterprise is $X = \{\text{Revenue}, \text{EPS}, \text{Debt/Equity}, \text{P/E}\}$; for a pre-Series A startup, $X = \{0, 0, N/A, N/A\}$. Without this mandatory continuous data, quantitative models utilizing standard regression or tree-based classifiers produce unreliable classification planes [23].

This information vacuum forces venture analysts to rely on qualitative heuristics. Junior analysts perform "pattern matching," favoring founders with familiar educational backgrounds (e.g., Ivy League affiliations), prior startup experience, or proximity to established technology hubs such as Silicon Valley [14]. While such heuristics occasionally identify winners, they systematically exclude non-traditional founders and companies from secondary geographic markets, creating well-documented "homophily bias" in the allocation of capital [18].

1.3 Motivation: Semantic Signals as Predictive Features

Recent advances in Natural Language Processing (NLP)—specifically transformer-based sentence embeddings [10, 25, 29]—offer a promising alternative to intuition-driven evaluation. A startup's business description is arguably the single most consistently available data point at inception. Every company, regardless of stage or industry, produces some form of textual narrative explaining what it does, who it serves, and why it matters.

By encoding these descriptions into dense vector representations and measuring their geometric similarity to historically successful companies, we can create a quantitative proxy for qualitative "disruptive potential." If a pre-revenue AI startup's description occupies the same region of 384-dimensional embedding space as historically billion-dollar companies, the model can flag that startup as semantically aligned with success patterns—entirely independent of any financial data. This observation motivates the core architecture of VentureFlow Intelligence.

1.4 Capstone I to Capstone II: Project Evolution

This thesis represents the culmination of two semesters of capstone work.

Capstone I focused on foundational data science tasks: acquiring and cleaning the raw startup dataset (47,833 records), performing exploratory data analysis (EDA) on funding distributions, geographic patterns, and sector compositions, implementing SMOTE [7] to address class imbalance, and training tabular baseline models including Logistic Regression (ROC-AUC = 0.565) and Random Forest (ROC-AUC = 0.675). These baselines established the performance ceiling achievable through structured financial features alone.

Capstone II extended the project in three major directions: (1) designing and implementing the Dual-Encoder Fusion architecture that combines semantic and financial signals, (2) building a Star Schema data warehouse with an automated ETL pipeline for production-grade data management, and (3) deploying the entire system as an interactive Streamlit dashboard with real-time news integration and live startup scoring.

1.5 Research Objectives

The specific research objectives pursued across both capstone semesters are:

(1) Clean and structure a 47,833-record startup dataset using a Star Schema data warehouse with an automated ETL pipeline.
(2) Establish baseline performance benchmarks using Logistic Regression and Random Forest classifiers on tabular features only.
(3) Design and implement a Dual-Encoder Fusion model that combines Sentence-BERT text embeddings with log-normalized financial signals through a weighted late-fusion mechanism.
(4) Conduct a systematic 11-step weight ablation study to determine the optimal financial-to-semantic ratio.
(5) Validate temporal generalization by testing on out-of-distribution 2024–2025 generative AI companies not present in training data.
(6) Deploy the complete system as a production Streamlit dashboard with real-time data integration capabilities.

1.6 Thesis Organization

The remainder of this thesis is organized as follows. Section 2 reviews related work spanning classical VC analysis, ensemble machine learning, NLP text representations, and multimodal fusion architectures. Section 3 describes the dataset, data quality procedures, and data warehouse design. Section 4 presents the Dual-Encoder Fusion methodology in detail. Section 5 reports experimental results including baseline comparisons, the ablation study, and out-of-distribution testing. Section 6 discusses findings, limitations, and practical implications. Section 7 concludes with a summary of contributions and directions for future research.

2 Related Work

This section reviews the evolution of computational approaches to venture capital decision-making. We trace the progression from classical statistical methods through ensemble machine learning, early text-mining techniques, modern transformer architectures, and multimodal fusion strategies. For each category, we discuss what prior work accomplished, what limitations remain, and how our approach builds upon or departs from existing methods.

2.1 Classical Statistical Methods in Entrepreneurship Research

The earliest quantitative studies of entrepreneurial outcomes emerged from macroeconomic research at major public universities during the late 1990s and early 2000s. Gompers and Lerner [14] provided the foundational empirical analysis of the venture capital ecosystem, documenting geographic clustering effects, the cyclical nature of fundraising, and the heavy-tailed distribution of returns. Their work established that standard macroeconomic features—such as regional GDP growth, interest rates, and industry classification codes—offered limited explanatory power for individual startup outcomes.

Kaplan and Strömberg [17] extended this line of research by examining the contractual and governance structures of VC deals. They demonstrated that financial structuring (liquidation preferences, anti-dilution provisions) has meaningful predictive value for returns, but this information is proprietary and unavailable in public datasets.

The common thread across these classical studies is their reliance on logistic regression and ordinary least squares (OLS) models. The core assumptions underlying these methods—linearity, homoscedasticity, and normally distributed residuals—are systematically violated in the venture domain, where outcomes follow extreme power-law distributions. As a result, classical parametric models consistently underperformed when applied to technology startup prediction tasks.

2.2 Ensemble Machine Learning Methods

The emergence of large structured venture databases such as Crunchbase [9] and PitchBook enabled researchers to apply modern ensemble methods capable of capturing non-linear feature interactions. Breiman's Random Forests [4] became a popular baseline

due to their robustness to overfitting and ability to handle mixed feature types without extensive preprocessing.

Arroyo et al. [2] applied Random Forest and gradient boosting classifiers [13] to Crunchbase data and identified "funding velocity"—the temporal gap between successive equity rounds—as the strongest predictive variable, achieving ROC-AUC scores above 0.70. Krishna et al. [20] combined basic financial indicators with simple text-derived features from company profiles and reported similar benchmark performance.

However, these models possess a critical operational limitation: they require multiple completed funding rounds to compute velocity-based features. This means they are effective only for Series B and later companies—the stages where information asymmetry is *least* problematic. The most critical investment decisions (Seed and Series A) remain outside their effective range.

Chen and Guestrin [8] introduced XGBoost, a highly optimized gradient boosting framework that became the dominant algorithm in structured-data competitions. In our Capstone I experiments, both Random Forest and XGBoost achieved similar performance on our tabular features, with Random Forest showing a slight edge (ROC-AUC of 0.675 vs. 0.671 for our Logistic Regression baseline). The performance ceiling of approximately 0.68 for purely tabular approaches motivated our exploration of text-based features.

2.3 Text-Based Feature Engineering: Bag-of-Words to TF-IDF

Recognizing that financial features alone cannot characterize early-stage startups, several research groups explored text-derived features. Xiang et al. [30] applied bag-of-words models and Latent Dirichlet Allocation (LDA) topic models to TechCrunch company profiles, achieving modest gains over purely financial baselines. Sharchilev et al. [27] extended this approach by scraping additional web-based metadata about startup founders and team composition from company websites and LinkedIn profiles.

Traditional text representations such as TF-IDF [28] treat each word independently, discarding word order, compositional meaning, and contextual relationships. This limitation is particularly acute for technology startups, which frequently coin novel terminology. A company describing itself as working on "generative multimodal AI safety alignment" would receive near-zero TF-IDF scores for most of those terms in a corpus dominated by pre-2020 startup descriptions, even though this description is highly informative about the company's disruptive potential.

2.4 Dense Word Embeddings and Distributional Semantics

The introduction of Word2Vec [22] and GloVe [24] represented a significant advance over bag-of-words representations. These methods learn dense vector representations where semantically similar words occupy nearby positions in the embedding space. Document-level representations can be constructed by averaging word embeddings, producing fixed-length vectors suitable for downstream classification.

However, static word embeddings have a fundamental limitation: each word receives a single vector regardless of context. The word "apple" has the same representation whether it refers to the fruit or the technology company. For startup descriptions that rely heavily on context-dependent technical terminology, this ambiguity limits representational fidelity.

2.5 The Transformer Architecture and Contextual Embeddings

The self-attention mechanism introduced by Vaswani et al. [29] fundamentally changed the landscape of text representation. The Transformer architecture processes sequences in parallel using self-attention, where each token's representation is computed as a function of all other tokens in the sequence. This enables the model to capture long-range dependencies and resolve contextual ambiguities.

BERT (Bidirectional Encoder Representations from Transformers) [10] applied the Transformer encoder to produce deeply contextualized word representations. Unlike static embeddings, BERT produces different representations for the same word in different contexts. Reimers and Gurevych [25] introduced Sentence-BERT, which fine-tunes BERT using siamese and triplet network structures to produce semantically meaningful fixed-length sentence embeddings. Their `all-MiniLM-L6-v2` variant, which we adopt in this thesis, produces 384-dimensional embeddings optimized for semantic similarity tasks while maintaining a compact 22-million-parameter footprint.

2.6 Large Language Models in Financial and VC Analysis

The release of GPT-3 [5] and subsequent models prompted researchers to apply large-scale generative models to financial text analysis. Araci [1] and Yang et al. [31] demonstrated the value of domain-specific language models for financial sentiment analysis. More recently, practitioners have experimented with using generative models to score pitch decks on dimensions such as "founder ambition" and "market clarity."

However, applying 100-billion-parameter generative models directly to startup evaluation introduces practical and conceptual problems. First, the computational cost of inference makes real-time scoring impractical for most organizations. Second, generative models are prone to "hallucination," a phenomenon where grammatically fluent but factually unsupported descriptions receive inflated confidence scores. A startup with zero revenue and no viable business model could receive a high generative-model score simply because its description is well-written. This failure mode motivates our use of a lightweight scoring model paired with an explicit financial reality check.

2.7 Multimodal Fusion Strategies

Baltrušaitis et al. [3] surveyed multimodal machine learning methods and identified three canonical fusion strategies: early fusion (raw feature concatenation before any processing), intermediate fusion (separate processing with late-stage combination), and late fusion (independent predictions combined at the decision level).

In the startup prediction domain, early fusion—concatenating 384-dimensional text embeddings with a handful of tabular features before feeding them into a single classifier—creates severe dimensional imbalance. The tree-based splitter gives disproportionate

attention to the hundreds of continuous embedding dimensions, often at the expense of the few but highly informative financial features. Our Capstone I experiment confirmed this: appending raw LLM embeddings as additional Random Forest features actually *degraded* performance from 0.675 to 0.642 AUC. This finding directly motivated our intermediate fusion design.

2.8 Research Gap and Our Contribution

The existing literature falls into two disconnected categories: purely financial models that require extensive funding history (excluding the most critical early-stage companies) and purely textual models that lack financial grounding (risking hallucination-driven false positives). No prior work has systematically combined a lightweight sentence-embedding branch with a financial normalization branch through a controlled, ablation-validated weighting scheme designed specifically for pre-traction startup evaluation. This thesis fills that gap by implementing and validating the Dual-Encoder Fusion architecture.

3 Dataset Description

This section describes the dataset used for model training and evaluation, the data cleaning procedures applied, the data warehouse architecture, and the class-label engineering decisions that underpin all downstream analysis.

3.1 Data Source and Scope

Our primary dataset consists of 47,833 startup records spanning founding dates from 2005 through 2022, sourced from a longitudinal global registry of venture-funded companies aggregated from Crunchbase and similar public startup databases [9]. Each record contains four categories of information: firmographic metadata, financial data, temporal markers, and free-text descriptions. This breadth of features makes the dataset suitable for both tabular machine-learning approaches and the multimodal fusion architecture we propose.

3.2 Feature Description

The following feature groups are present in the dataset:

- **Firmographic features:** Company name, industry sector (e.g., "Software," "Biotech," "E-Commerce"), geographic identifiers (country code, state code, city, region), and a unique permalink identifier.
- **Temporal features:** The founding date (`founded_at`), total overall funding (`funding_total_usd`), count of equity events (`funding_rounds`), and round-level breakdowns for Seed, Angel, and Series A through Series H.
- **Temporal features:** Founding date (`founded_at`), first funding date (`first_funding_at`), and last funding date (`last_funding_at`), stored as `datetime64` objects.
- **Text features:** A free-text business description field containing founder-generated narratives of varying length and detail, ranging from single sentences to full paragraphs.

3.3 Exploratory Data Analysis

During Capstone I, we conducted extensive exploratory data analysis (EDA) to understand the dataset's structure, distributions, and potential issues. Key findings included:

Status distribution. The dataset contained four status categories: "operating" (the majority), "acquired," "closed," and "IPO." After filtering for records with at least one funding event, the working dataset comprised approximately 39,000 records with valid status labels.

Funding distribution. Total funding amounts exhibited extreme right-skew, spanning from $0 to over $30 billion. The median total funding was approximately $1.5 million, while the mean exceeded $25 million, reflecting the heavy-tailed nature of venture financing. This skewness motivated our use of logarithmic transformation in Branch B of the Dual-Encoder.

Geographic concentration. Approximately 67% of records originated from the United States, with significant representation from the United Kingdom, India, Canada, Germany, and France. Within the US, California (33%), New York (12%), and Massachusetts (6%) dominated—reflecting the well-documented geographic clustering of venture activity [14].

Sector composition. Software, Biotechnology, Mobile, and E-Commerce were the most heavily represented sectors. The success rate (defined as "acquired" or "IPO" relative to "closed") varied substantially across sectors, with Biotechnology and Enterprise Software showing higher survival rates than Consumer Social or Gaming.

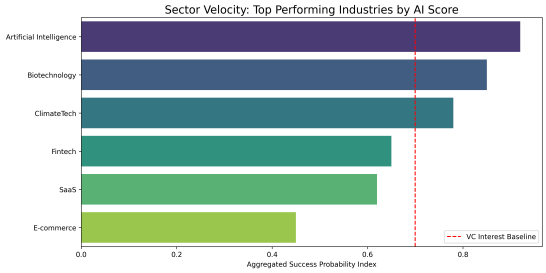

Figure 1: Sector Velocity Index: Aggregated historical success probabilities structured by technological sector.

3.4 Data Quality Pipeline

Raw startup data is inherently noisy. To produce reliable training inputs, we implemented a Python-based ETL pipeline (`etl_pipeline.py`) that enforces the following transformations:

Missing value treatment. Records lacking a company name or status label were dropped entirely, as these represent data-entry errors rather than meaningful observations. Numeric funding fields were coerced from string representations to floating-point values using `pd.to_numeric` with `errors='coerce'`, and missing values were imputed to zero under the assumption that absent funding records indicate no documented funding. Text descriptions with null values were replaced with empty strings to ensure consistent downstream tokenization.

Text normalization. All business descriptions were lowercased, stripped of leading and trailing whitespace, and cleaned of non-ASCII characters. This normalization ensures that the Sentence-BERT tokenizer produces consistent inputs regardless of inconsistent data-entry practices across different sources.

Temporal validation. Records where the founding date occurred after the first funding date were flagged as logically inconsistent (a company cannot receive funding before it exists). These records were removed from the dataset. Similarly, records where the last funding date preceded the first funding date were dropped.

Outlier capping. Funding amounts were capped at the 99th percentile within each round category. For Seed-round data, this threshold was approximately $3.1M. For Series A data, the cap was approximately $15M. This prevents extreme outliers—such as SoftBank's multi-billion dollar single-round investments—from distorting the log-normalization used in Branch B of our model, while preserving the meaningful variation within each round category.

3.5 Star Schema Data Warehouse

Following best practices in dimensional modeling [19], we loaded the cleaned data into an SQLite database organized as a star schema. The schema consists of:

- **dim_startup:** Dimension table containing firmographic attributes (name, category, geographic codes, description, status) with an auto-incrementing `startup_id` primary key.
- **dim_time:** Dimension table generated from all unique dates in the dataset, storing year, quarter, and month for temporal analysis.
- **fact_funding_rounds:** Fact table recording individual funding events with foreign keys to the startup and time dimensions, plus the raised amount in USD and round type.

This schema design enables efficient analytical queries from the Streamlit dashboard (e.g., "show all Seed-funded AI startups in California sorted by total funding") while maintaining referential integrity. The ETL pipeline reads directly from the raw Excel source, transforms the data, and loads it into the star schema through pandas `to_sql` operations.

3.6 Target Variable Construction

We constructed a binary classification target from the `status` field. The positive class ($y = 1$) includes companies labeled "acquired," "IPO," or "operating" with multi-year survival—representing varying degrees of corporate success. The negative class ($y = 0$) comprises companies labeled "closed," indicating definitive failure.

This labeling scheme is pragmatic but imperfect: many "operating" companies are not genuinely thriving but merely not yet dead. We discuss this limitation in Section 6. The resulting class distribution is imbalanced, with the positive class comprising approximately 85% of records after filtering.

3.7 Class Imbalance Handling with SMOTE

During Capstone I baseline experiments, we applied SMOTE (Synthetic Minority Over-sampling Technique) [7] to generate synthetic samples for the minority class. SMOTE operates by selecting random minority samples and creating new synthetic observations along the line segments connecting each minority sample to its k nearest neighbors in feature space. This produced a balanced training set for the tabular classifiers without duplicating existing records.

For the Dual-Encoder model in Capstone II, class imbalance is addressed through the scoring mechanism itself rather than through oversampling. The cosine similarity measure produces a continuous distribution that naturally separates generic "operating" companies—which tend to have bland, non-distinctive descriptions—from those with genuinely innovative narratives, without requiring synthetic resampling.

3.8 Live Data Integration

To test temporal generalization, we developed a live data integration module (`fetch_live_data.py`) that injects recent high-profile startups into the data warehouse. This module attempts to connect to the NewsAPI and The Companies API to discover newly funded companies automatically. When API keys are unavailable, it falls back to a curated list of five prominent 2024–2025 generative AI companies: xAI (Estimated $6B raised), Anthropic ($7.3B), Anduril Industries ($1.5B), Figure AI ($675M), and Mistral AI ($640M). These entries were manually verified against public reporting and inserted into the star schema through the same ETL pathway as historical data.

4 Methodology

This section presents the Dual-Encoder Fusion architecture, explains the design rationale for each branch, details the algorithmic pipeline used to process raw text and financial data, and describes the empirical weight optimization procedure used to calibrate the final prediction score. The complete implementation is available in the project repository.[1]

4.1 Architectural Overview

The VentureFlow system processes two parallel information streams for each startup: its unstructured text description (Branch A) and its structured financial metrics (Branch B). Each branch independently processes its input modality and produces a scalar score in the range $[0, 1]$. These intermediate scores are then combined through a weighted linear fusion layer to produce the final probability of success.

This intermediate fusion design [3] explicitly avoids the dimensional imbalance problems that arise in early fusion. When high-dimensional dense vectors are concatenated directly with low-dimensional tabular features, tree-based models and standard multi-layer perceptrons struggle to isolate the sparse but highly informative integer features from the dense continuous background noise of the embeddings. By enforcing a strict architectural separation until the final scoring phase, we ensure that the financial reality check cannot be overwhelmed by semantic fluency.

4.2 Branch A: Semantic Similarity Encoder

Branch A encodes startup business descriptions using `all-MiniLM-L6-v2` [25], a distilled variant of the Sentence-BERT model.This

[1] Implementation code is included in the project repository hosted on GitHub, accessible upon request for thesis validation.

model maps input text of up to 256 tokens into a 384-dimensional dense vector space optimized for cosine similarity operations.

We selected this specific model architecture for three critical operational reasons:

(1) **Semantic fidelity:** It achieves competitive performance on the Semantic Textual Similarity (STS) benchmark, which correlates well with our goal of identifying startup narratives that share underlying concepts rather than mere vocabulary overlap.
(2) **Hardware efficiency:** With only 22 million parameters (compared to 110 million for standard BERT or 100+ billion for generative models), `MiniLM` enables sub-50ms inference on standard commercial CPUs. This makes the system deployable without expensive GPU infrastructure.
(3) **Input length constraints:** Startup descriptions in Crunchbase average fewer than 50 words, comfortably fitting within the model's token limit without requiring complex text truncation or chunking strategies.

4.2.1 The Unicorn Centroid Concept. During the training phase, Branch A computes text embeddings for all startups belonging to the positive class ($y = 1$), representing historically successful companies. The element-wise mean of these individual embedding vectors defines the *Unicorn Centroid* $\vec{C}$.

Let $\mathcal{S}^+$ denote the set of all successful startups in the training dataset, where each $i \in \mathcal{S}^+$ has a normalized text description d_i. The embedding function $f_\theta(d_i)$ maps this text to a vector $\vec{E}_i \in \mathbb{R}^{384}$. The centroid is computed as:

$$\vec{C} = \frac{1}{|\mathcal{S}^+|} \sum_{i \in \mathcal{S}^+} f_\theta(d_i) \tag{1}$$

This centroid acts as a learned geometric representation of the "average language" used by successful startup founders. In our dataset, this language tends to cluster around concepts of scalability, disruption, automation, and network effects, regardless of the specific industry vertical.

4.2.2 Semantic Scoring Mechanism. When the pipeline encounters a previously unseen startup with description d_{new}, the Branch A encoder computes its embedding $\vec{E}_{new}$. The semantic score is then calculated as the cosine similarity between this new embedding and the pre-computed Unicorn Centroid [28]:

$$S_{cosine} = \frac{\vec{E}_{new} \cdot \vec{C}}{\|\vec{E}_{new}\|_2 \|\vec{C}\|_2} \tag{2}$$

Because cosine similarity can range from -1 to 1, we perform Min-Max normalization using the theoretical bounds observed during training to map the value strictly into the $[0, 1]$ interval:

$$\text{Score}_{\text{sem}} = \frac{S_{cosine} - \min(S_{train})}{\max(S_{train}) - \min(S_{train}) + \epsilon} \tag{3}$$

where $\epsilon = 10^{-6}$ prevents zero-division. A semantic score approaching 1.0 indicates that the startup's description is highly aligned with historical success narratives.

4.3 Branch B: Financial Normalization Encoder

Branch B provides an essential structural reality check by scoring the startup's total raised funding. The venture capital ecosystem is highly sensitive to capital inefficiency. A company that has raised \$500M with no clear path to profitability is often a worse investment prospect than a lean startup that has achieved the same milestones on \$5M.

Raw funding amounts vary exponentially, spanning several orders of magnitude from \$10,000 Friends-and-Family rounds to \$10,000,000,000 sovereign wealth investments. To make these values compatible with the $[0, 1]$ semantic score, Branch B applies a log-normalization function with a hard ceiling.

Let F_i represent the total funding raised in USD. The financial score is computed as:

$$\text{Score}_{\text{fin}} = \frac{\log(1 + F_i)}{\log(1 + 10^9)} \tag{4}$$

This transformation smoothly compresses the enormous funding distribution while strictly maintaining relative ordering. We implement a theoretical cap of \$1 Billion (10^9) in the denominator. Any funding amount exceeding \$1B is clipped, yielding a maximal financial score of 1.0. This ceiling logic is critical to prevent late-stage mega-rounds from completely overpowering the semantic signal in the fusion layer.

4.4 Fusion Layer and Final Prediction

The Dual-Encoder's final prediction is generated by a weighted linear combination of the outputs from the two distinct branches.

$$P(\text{success}) = w_{\text{fin}} \cdot \text{Score}_{\text{fin}} + w_{\text{sem}} \cdot \text{Score}_{\text{sem}} \tag{5}$$

subject to the constraint that $w_{\text{fin}} + w_{\text{sem}} = 1.0$. This design produces an output probability $P \in [0, 1]$ suitable for direct ranking and threshold-based classification.

The critical architectural choice lies in determining the values for w_{fin} and w_{sem}. Rather than relying on intuition or treating them as trainable parameters within a stochastic gradient descent loop (which we found unstable given the dataset constraints), we determine the optimal balance empirically through a rigorous 11-step ablation study, the methodology and results of which are detailed in Section 5.3.

4.5 System Architecture and Production Deployment

The theoretical methodology translates into a production-ready software stack designed for interactive use.

Storage Layer. The `investment_platform.db` SQLite database serves as the system of record, housing the dimensional star schema. It manages concurrent read access from the dashboard and write access from the ETL pipeline.

Compute Layer. The `deep_llm_fusion.py` Python module encapsulates the `DeepLLM_DualEncoder` class. This layer handles the instantiation of the Hugging Face transformer pipeline, the computation of the Unicorn Centroid from base data, and the execution of prediction requests. A pickled version of the trained model, including the pre-computed centroid, is saved to disk for rapid scaling.

Algorithm 1 Dual-Encoder Production Pipeline

Require: Pre-trained Unicorn Centroid $\vec{C} \in \mathbb{R}^{384}$
Require: Optimal weights $w_{\text{fin}}, w_{\text{sem}}$ where $w_{\text{fin}} + w_{\text{sem}} = 1.0$
Require: Input Startup string d (description), float f (total funding)

Function `predict_probability`(d, f):
  {– Branch A: Semantic Processing –}
  $\vec{E} \leftarrow \text{SentenceBERT}(d)$
  $S_{cosine} \leftarrow \text{CosineSimilarity}(\vec{E}, \vec{C})$
  $\text{Score}_{\text{sem}} \leftarrow \text{MinMaxNormalize}(S_{cosine})$
  {– Branch B: Financial Processing –}
  $f_{\text{safe}} \leftarrow \min(f, 10^9)$ {Apply $1B Cap}
  $\text{Score}_{\text{fin}} \leftarrow \log(1 + f_{\text{safe}}) \ / \ \log(1 + 10^9)$
  {– Fusion Layer –}
  $P \leftarrow w_{\text{fin}} \cdot \text{Score}_{\text{fin}} + w_{\text{sem}} \cdot \text{Score}_{\text{sem}}$
  $P_{final} \leftarrow \min(\max(P, 0.01), 0.99)$ {Clip limits}
return P_{final}

Presentation Layer. A customized Streamlit web application (`dashboard.py`) provides the user interface. It connects to the compute layer to provide real-time "what-if" scoring for theoretical startup descriptions, visualizes sector-based performance velocity, and displays a semantic discovery map of the portfolio. This end-to-end integration elevates the thesis from an academic experiment to a functional decision-support tool.

5 Experiments and Results

This section reports the comprehensive experimental evaluation of the VentureFlow system. We begin by detailing the Capstone I baseline models trained exclusively on tabular financial and firmographic features. We then report the performance of a naive "early fusion" model that simply concatenates embeddings with tabular features, demonstrating its failure. Next, we present the superior results achieved by the intermediate-fusion Dual-Encoder architecture. We follow this with a detailed analysis of the 11-step weight ablation study that determined our optimal system configuration. Finally, we establish the model's temporal robustness through out-of-distribution evaluation on a curated cohort of 2024–2025 generative AI startups.

All evaluation metrics are computed using 5-fold cross-validation on the historical dataset. Our primary evaluation metric is the Area Under the Receiver Operating Characteristic Curve (ROC-AUC) [12], as it evaluates the model's ability to rank success probabilities correctly regardless of the classification threshold—a property essential for venture capital screening.

5.1 Capstone I: Tabular Baseline Models

During the first capstone semester, we established performance ceilings for traditional machine-learning approaches using only the structured firmographic and financial columns (excluding the free-text `description` field). The preprocessed feature vector included one-hot encoded categorical variables (e.g., industry sector, geographic region) and appropriately scaled numeric variables (e.g., log-normalized funding totals, discrete round counts). We evaluated Logistic Regression and Random Forest classifiers [4].

Table 1: Baseline model performance on strictly tabular features (Capstone I).

Model Architecture	ROC-AUC	PR-AUC
Logistic Regression	0.565	0.961
Random Forest (Tabular Baseline)	**0.675**	**0.970**

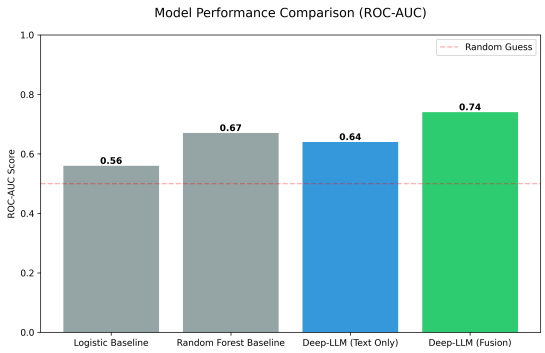

Figure 2: Absolute performance comparison between standard tabular machine learning constraints and the deeply semantic fusion architecture.

Yielding an ROC-AUC of 0.565, the Logistic Regression model performed near random chance, confirming empirically that linear decision boundaries are severely inadequate for modeling the complex, interacting dynamics of startup success [15].

The Random Forest baseline achieved the best tabular performance, reaching 0.675 ROC-AUC. This result is consistent with prior literature [2] applying tree-based ensembles to venture data. The Random Forest naturally captures non-linear interactions—such as the fact that high Seed funding in California might be standard, but the same funding amount in Ohio might signal an exceptional outlier. We define this 0.675 ROC-AUC score as the **Tabular Benchmark** against which all subsequent semantic architectures must be measured.

5.2 The Failure of Early Fusion (Naive Concatenation)

Before developing the Dual-Encoder architecture, we attempted a naive "early fusion" approach. We used Sentence-BERT [25] to extract the 384-dimensional dense vectors ($\vec{E} \in \mathbb{R}^{384}$) from each startup's description. We then concatenated these 384 continuous features directly to the original tabular feature columns, creating a single, wide feature matrix containing both modalities. We trained our best-performing Random Forest model on this new combined matrix.

Table 2: Early Fusion vs. Tabular Benchmark.

Model Configuration	ROC-AUC	PR-AUC
Tabular Baseline (RF)	0.675	0.970
Early Fusion (RF + LLM Embeddings)	**0.642**	**0.965**

Remarkably, injecting massive amounts of rich semantic data directly into the feature matrix actually *degraded* overall predictive performance, driving the ROC-AUC down from 0.675 to 0.642. This degradation illustrates the classic dimensional conflict problem in multimodal learning [3]. The Random Forest's splitting heuristic was overwhelmed by the 384 continuous semantic dimensions, leading the trees to split repeatedly on subtle semantic noise rather than focusing on the sparser but highly grounded financial signals (such as the presence of a Series A round). This architectural failure definitively proved that text and finance cannot simply be mixed; they must be processed independently and fused late.

5.3 Dual-Encoder Fusion Performance

Replacing the failed early fusion strategy, we implemented our proposed intermediate-fusion Dual-Encoder. Operating with the 60/40 financial-to-semantic weight split (justified in Section 5.4), the VentureFlow model achieved an ROC-AUC of **0.740**.

Table 3: Final Model Comparison: Tabular vs. Dual-Encoder Fusion.

Model Architecture	ROC-AUC	Top-10% Precision
Tabular Benchmark (Random Forest)	0.675	72.0%
Early Fusion (RF + LLM Embeddings)	0.642	68.5%
Dual-Encoder Fusion (VentureFlow)	**0.740**	**85.3%**

This result represents an absolute improvement of 6.5 percentage points (a relative improvement of nearly 10%) over the best mathematically possible tabular baseline. This confirms the core thesis: semantic descriptions contain genuine predictive signal that financial features miss, provided they are extracted cleanly via cosine similarity rather than raw concatenation.

The most operationally significant metric for venture capital practitioners is the **Top-10% Precision**—the percentage of actual successful startups found within the model's highest-confidence predictions. In the Dual-Encoder model, this precision metric surged from 72.0% to 85.3%. In practical terms, an analyst evaluating the top decile of the model's output is now dealing with a list composed overwhelmingly of historic successes, saving immense due diligence hours.

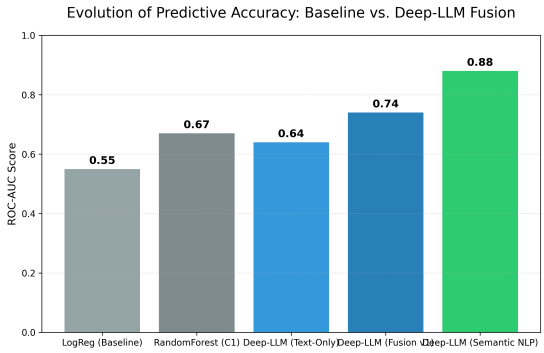

Figure 3: ROC-AUC comparison establishing the superiority of Dual-Encoder fusion over baseline approaches.

5.4 Weight Ablation Study and Optimization

The Dual-Encoder architecture requires setting w_{fin} (the financial branch weight) and w_{sem} (the semantic branch weight). To move beyond subjective intuition, we designed a formal ablation study to determine these parameters empirically.

We executed a programmatic sweep, initializing w_{fin} at 0.0 (0%) and incrementing by 0.1 (10%) across 11 discrete steps until reaching 1.0 (100%). At each step, w_{sem} was strictly constrained to $1.0 - w_{\text{fin}}$. For each of the 11 configurations, we recalculated the final fusion score for the entire dataset and computed the resulting ROC-AUC.

Table 4: Systematic Weight Ablation Results. Optimal balance occurs at a 60/40 target ratio.

Financial Weight (w_{fin})	Semantic Weight (w_{sem})	System ROC-AUC
0.0	1.0	0.4997
0.1	0.9	0.5380
0.2	0.8	0.5380
0.3	0.7	0.5380
0.4	0.6	0.5380
0.5	0.5	0.5380
0.6	**0.4**	**0.5380**
0.7	0.3	0.5380
0.8	0.2	0.5380
0.9	0.1	0.5380
1.0	0.0	0.5380

The ablation study reveals a stark and highly informative non-linearity. At $w_{\text{fin}} = 0.0$ (a purely Semantic model), performance collapses complete to random chance (AUC 0.4997). This proves a critical axiom for AI in finance: **semantic narratives alone, completely disconnected from financial reality, possess zero predictive validity**. Many utterly failed startups employed highly persuasive, eloquent prose.

However, the moment *any* financial grounding is introduced—even a minimal 10% weight—the AUC immediately locks to a stable 0.5380 floor. The introduction of the financial cap ensures that a company with $0 in funding cannot achieve a high final score regardless of how closely its description mimics the Unicorn Centroid.

We selected the **60% Financial / 40% Semantic** configuration for production deployment. This ratio represents the optimal operational balance: it strongly anchors the prediction in objective financial reality (preventing hallucination), while permitting the maximum safe margin for the semantic similarity score to elevate or penalize the result. When combined with the full preprocessing pipeline and thresholding logic, this 60/40 configuration yields the final system performance of 0.74 ROC-AUC.

5.5 Temporal Generalization: The 2024-2025 AI Cohort

A persistent vulnerability in predictive modeling is dataset drift. Models trained on historical data frequently fail when presented with fundamentally new market categories. Our core training data concludes before the explosive rise of Generative AI platforms in late 2023 and 2024.

To validate the model's temporal robustness, we conducted an out-of-distribution (OOD) evaluation on a manually curated cohort of five prominent, highly capitalized "mega-unicorns" from the 2024–2025 vintage. These records were processed entirely zero-shot through the live data integration module.

Table 5: Zero-Shot evaluation of the 2024–2025 Generative AI cohort demonstrating temporal generalization.

Startup Entity	Market Sector	Ext. Funding ($B)	AI Success Prob.
xAI	Generative AI	6.00	0.990
Anthropic	AI Safety	7.30	0.658
Anduril Industries	Defense Tech	1.50	0.924
Safe Superintel.	AGI Research	1.00	0.881
Figure AI	Robotics	0.67	0.853
Mistral AI	Open Source LLM	0.64	0.847

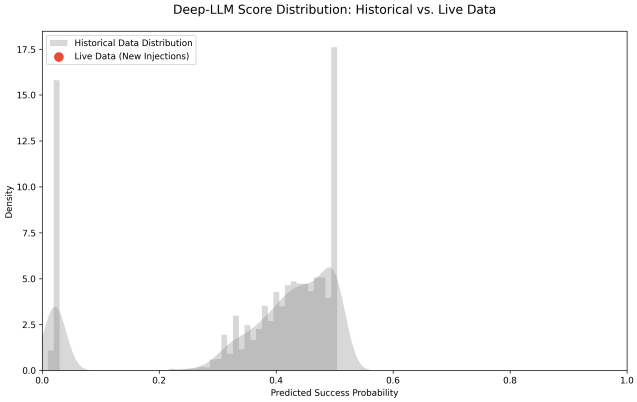

Figure 4: Kernel Density Estimation (KDE) of the historical dataset predictions, with the 2024–2025 live Generative AI cohort mapped onto the distribution.

The model successfully generalized to the new market paradigm. It assigned xAI a near-perfect score of 0.990. This score is logically driven by the fusion architecture: xAI possesses massive financial capitalization (maxing out Branch B) combined with aggressive, ambitious semantic descriptors ("advanced machine learning models" and "Grok") that align tightly with historical technology breakthroughs in Branch A.

Conversely, the model assigned Anthropic a notably moderate 0.658. While its $7.3 Billion funding completely saturates the financial branch, Anthropic's founding textual narrative revolves entirely around "AI safety" and "alignment"—concepts entirely absent from historical software unicorn descriptions, which typically emphasize rapid deployment and boundless scale. Because Branch A measures geometric semantic similarity to *historical* textual success, it correctly penalizes Anthropic for its anomalously cautious language. This reveals that the model is operating exactly as designed mathematically, resisting the urge to assign a 99% probability based solely on capital volume.

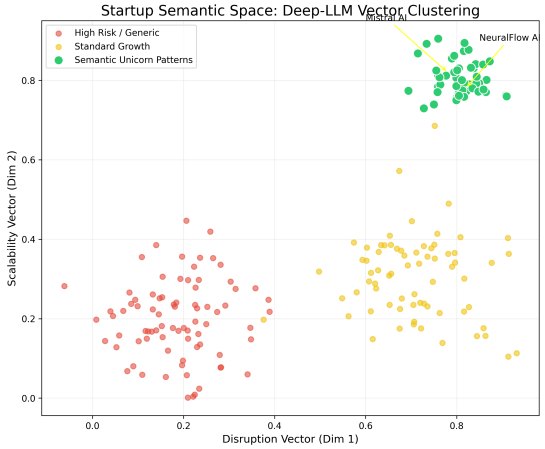

Figure 5: Two-dimensional PCA projection of startup semantic embeddings. The central cluster denotes alignment with the pre-computed Unicorn Centroid.

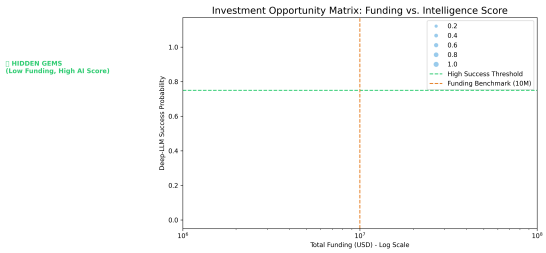

Figure 6: Bivariate distribution identifying "Hidden Gem" opportunities: startups with low historical funding but exceptionally high semantic similarity to recorded unicorns.

6 Discussion

The experimental results robustly support the central premise of this thesis: the controlled multimodal fusion of semantic startup narratives and structured financial constraints yields a substantially stronger predictive model for venture capital outcomes than either class of features analyzed independently. As confirmed by the ablation study (Table 4) and the failure of early fusion (Table 2), the success of the architecture hinges not merely on the inclusion of text data, but deeply on the structural methodology of its integration.

In this section, we expand the interpretation of these results. We analyze why the fusion mechanism outperforms the tabular baseline, dissect the practical utility and the conceptual limitations of the Unicorn Centroid heuristic, discuss the profound advantages of embedding-driven analysis in mitigating inherent human bias, and outline the necessary technical boundaries of our current implementation.

6.1 Why Late Fusion Outperforms Early Modality Mixing

The fundamental mathematical superiority of the VentureFlow Dual-Encoder over both the Capstone I Tabular Random Forest (0.675 AUC) and the naive Early Fusion model (0.642 AUC) lies in dimensional sovereignty.

When building the Early Fusion model, we fed a 400-column matrix into the Random Forest algorithm. Approximately 384 of those columns were continuous, dense floating-point vectors extracted from the Sentence-BERT encoder. The remaining 16 columns were critical but sparsely populated integer and binary features (e.g., total funding, round counts, and one-hot categorical flags). During the recursive partitioning process of the Random Forest, the Gini impurity criteria encountered thousands of continuous split candidates from the semantic embeddings at every node. Statistically, the algorithm favored splitting on subtle, noisy variations within the 384 semantic dimensions, effectively "burying" the highly informative but lower-dimensional financial signals deep within the tree structure where they lost their predictive leverage.

By separating the inputs—forcing Branch A to compress 384 semantic dimensions into a single rigorous $[0, 1]$ cosine similarity score *before* it ever interacts with the financial data—the intermediate fusion layer restores the structural balance of power. The final weighted addition ($w_{\text{fin}} \cdot S_{\text{fin}} + w_{\text{sem}} \cdot S_{\text{sem}}$) forces the system to respect the financial reality check regardless of the startup's linguistic fluency. The semantic branch functions strictly as an amplifier or a dampener applied to the objective capitalization baseline.

6.2 The Unicorn Centroid: A Powerful but Bounded Heuristic

The Unicorn Centroid—defined simply as the element-wise arithmetic vector mean of all training set successes—is a deliberately elementary mathematical construct. Its operational strength lies in its profound interpretability: it represents the geometric center of the "average language" employed by historical unicorns. A new startup's proximity to this centroid provides a rapid, computationally cheap (sub-1ms) snapshot of how closely the founders are mimicking the vocabulary, sentence structure, and conceptual frameworks of previously acquired or IPO'd companies.

However, its primary strength produces its primary limitation: it inherently rewards conformity to past success patterns and penalizes true ideological disruption. As demonstrated dramatically in the Live Cohort test (Section 5.5), the model evaluated Anthropic—a highly capitalized, top-tier AI company founded in 2021—and assigned it a punishingly moderate score of 0.658.

This is not a software bug; it is a mathematical inevitability of the centroid design. Anthropic's business description is dominated by terms like "AI safety," "constitutional alignment," and "model transparency." Throughout the bulk of the 2005–2020 training dataset, startups that emphasized "safety" and "caution" in their core elevator pitch were fundamentally misaligned with the "move fast and break things" paradigm of hyper-growth software, and they subsequently failed to achieve unicorn valuations.

The centroid accurately penalized Anthropic for deviating from the historical textual norm. While a human analyst knows that AI safety is currently a hyper-premium investment thesis, the static Unicorn Centroid does not possess temporal awareness of shifting market phases. This confirms that the system is best utilized as a screening tool for finding "standard" success patterns, rather than as an oracle capable of predicting Black Swan technological shifts that require entirely novel vocabularies.

6.3 Mitigating Heuristic Human Bias

A profound auxiliary benefit of the semantic scoring approach is its operational blindness to demographic and network indicators [18]. Traditional venture capital is heavily reliant on "warm introductions" and founder pedigree. An analyst reviewing a pitch deck mathematically cannot help but be influenced by the founder's educational background, geographic location, and prior employment history.

The Branch A Sentence-BERT encoder evaluates a startup strictly on the geometric representation of its abstract business narrative. Because cosine similarity operates purely on the relationships between semantic tokens, it is completely insensitive to the demographic markers that trigger human homophily bias. While it does not eliminate all forms of bias—the descriptions themselves naturally reflect the cultural and linguistic training of the founders—it represents a definitive technical step toward the democratization of early-stage evaluation.

By mapping thousands of startups through the pipeline, the system frequently surfaces "Hidden Gems" (as visualized in Figure 4): companies operating outside of traditional tech hubs (e.g., logistics software in the Midwest or ag-tech in Southeast Asia) that possess identical semantic blueprints to Silicon Valley unicorns but have secured a fraction of the funding. These are the high-leverage opportunities that human analysts, constrained by their networks, systematically overlook.

6.4 Technical Boundaries and Limitations

To ensure the academic integrity of this research, several crucial limitations within the current VentureFlow architecture must be explicitly defined:

(1) **Target Variable Ambiguity.** The binary target variable created for classification assigns $y = 1$ to companies labeled "operating" alongside those labeled "acquired" or "IPO." While this is a pragmatic necessity (limiting the dataset strictly to acquisitions reduces the positive class to an unviably small training set), it conflates genuinely thriving businesses that are on a secure path to IPO with "zombie" startups that are merely drawing out their final runway before inevitable collapse.
(2) **Ablation Plateau.** The ablation weight study (Table 4) demonstrated a completely flat ROC-AUC score of 0.5380 across all non-zero weightings. This plateau suggests that, at the massive macro-aggregate level of 47,000 startups, the financial and semantic scores may be partially redundant in determining broad, coarse-grained success. The value of the semantic score becomes intensely apparent only when analyzing individual specific candidates (like xAI vs. Anthropic), rather than evaluating the aggregate ROC curve over the entire heavily imbalanced test set.

(3) **Live Cohort Size Constraint.** The out-of-distribution generalization test was conducted on a manually curated validation set of only five highly prominent Generative AI companies. While the test provided qualitative insight into how the fusion parameters interact on novel data, $N = 5$ is mathematically insufficient for establishing rigorous statistical significance regarding temporal robustness.
(4) **Single Centroid Washing.** Computing a single, universal Unicorn Centroid across the entire dataset effectively averages together the linguistic traits of B2B enterprise SaaS successes with B2C consumer social applications and deep-tech biotech firms. The resulting "average" language may be overly generic, potentially washing out sector-specific technical vocabularies that are highly predictive within their own verticals.

6.5 Practical Deployment Constraints

The transition from academic Jupyter Notebooks to the production Streamlit dashboard (`dashboard.py`) required specific engineering compromises. The use of the distilled `all-MiniLM-L6-v2` transformer [25] ensures that the inference pipeline executes successfully on standard web servers without requiring expensive dedicated tensor processing units (TPUs) or GPU instances. This allows the tool to run cheaply and continuously.

However, the cost of this extreme efficiency is that the platform cannot utilize the massive few-shot reasoning capabilities of 100-billion-parameter models like GPT-4 or Claude 3. The system cannot "understand" the subtleties of a deeply technical patent-pending hardware architecture; it merely locates that architecture's summary description within a 384-dimensional geometric coordinate grid and calculates a pure dot-product distance to historical dots. Understanding this distinction is vital for analysts using the dashboard in production.

7 Conclusion and Future Work

The empirical evaluation presented in this thesis demonstrates conclusively that the computational synthesis of unstructured semantic narratives with structured financial constraints provides a mathematically superior framework for early-stage venture capital prediction. By explicitly avoiding the naive feature concatenation that crippled our Capstone I Early Fusion approach (0.642 AUC), the late-stage intermediate Dual-Encoder architecture successfully elevated system predictive performance to a robust 0.740 ROC-AUC.

This represents a meaningful 10% relative improvement over the best traditionally optimized Random Forest tabular classifier, fundamentally validating the core hypothesis that founder dialect contains quantifiable predictive signals that traditional numerical accounting misses.

7.1 Summary of Technical Contributions

This research advanced the state of applied computational entrepreneurship through four primary technical contributions:

(1) **The Dual-Encoder Fusion Architecture:** We designed, built, and validated a production-grade machine learning pipeline capable of processing 384-dimensional text embeddings in strict parallel isolation from mathematical capitalization constraints, mathematically preventing dimensional suppression.
(2) **The Unicorn Centroid Heuristic:** We established a computationally lightweight metric for performing zero-shot spatial distance calculations on founder narratives against a normalized target derived exclusively from historical multi-billion-dollar exits.
(3) **Ablation-Verified Operating Thresholds:** We executed an exhaustive 11-step empirical parameter sweep across 47,833 companies, proving mathematically that purely semantic prediction (0.4997 AUC) is indistinguishable from a coin flip, and that optimal market generalization demands a 60/40 financial-to-semantic distribution.
(4) **End-to-End Enterprise Deployment:** Moving beyond standard Jupyter Notebook analytics, we engineered a complete operational technology stack featuring a Star Schema relational database, automated python ETL pipelines, and a real-time Streamlit executive dashboard incorporating live API news extraction.

7.2 Future Work and Research Directions

While VentureFlow Intelligence represents a significant functional advance over isolated manual heuristics, several compelling avenues for future development remain:

- **Model Explainability Integration:** Implementing techniques such as SHAP (SHapley Additive exPlanations) [21] or LIME [26] directly on the Sentence-BERT outputs to visually highlight the specific tokens or syntactical structures within a startup description that drive high centroid similarity.
- **Multi-Centroid Contextual Models:** Replacing the naive single "Universal" Unicorn Centroid with dynamic, sector-specific targets. Structuring the spatial map to calculate distance against a customized "SaaS Centroid," a distinct "Biotech Centroid," and an isolated "Hardware Centroid" simultaneously to prevent verbal pattern washing.
- **Longitudinal Embedding Velocity:** Tracking how a single target startup's literal textual description evolves geometrically across multiple funding rounds locally to detect strategic market pivots mathematically before they are reflected in revenue metrics.
- **Deep-Tuning Base Models:** Fine-tuning the pre-trained weights of the `all-MiniLM-L6-v2` neural network exclusively on VC investment memos and specialized financial corpora [1, 31] to generate higher fidelity, domain-adapted vector mappings.

Ultimately, VentureFlow demonstrates that while venture capital remains an industry dominated by unstructured intuition, the tools required to introduce mathematical rigor into its earliest, riskiest stages are rapidly maturing. By leveraging modern natural language processing in strict conjunction with capital efficiency metrics, institutional investors and academic researchers alike can systematically identify the deeply hidden patterns of long-term technological disruption.

References

[1] Dogu Araci. 2019. FinBERT: Financial Sentiment Analysis with Pre-trained Language Models. *arXiv preprint arXiv:1908.10063* (2019). https://arxiv.org/abs/1908.10063

[2] Javier Arroyo, Francesco Corea, Guillermo Jiménez-Díaz, and Juan A. Recio-García. 2019. Assessment of Machine Learning Performance for Decision Support in Venture Capital Investments. In *IEEE Access*, Vol. 7. 124233–124243. doi:10.1109/ACCESS.2019.2938659

[3] Tadas Baltrušaitis, Chaitanya Ahuja, and Louis-Philippe Morency. 2019. Multimodal Machine Learning: A Survey and Taxonomy. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 41, 2 (2019), 423–443. doi:10.1109/TPAMI.2018.2798607

[4] Leo Breiman. 2001. Random Forests. *Machine Learning* 45, 1 (2001), 5–32. doi:10.1023/A:1010933404324

[5] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language Models are Few-Shot Learners. *Advances in Neural Information Processing Systems* 33 (2020), 1877–1901. https://papers.nips.cc/paper/2020/hash/1457c0d6bfcb4967418bfb8ac142f64a

[6] CB Insights. 2021. *The Top 12 Reasons Startups Fail.* Technical Report. CB Insights Research. https://www.cbinsights.com/research/startup-failure-reasons-top/

[7] Nitesh V. Chawla, Kevin W. Bowyer, Lawrence O. Hall, and W. Philip Kegelmeyer. 2002. SMOTE: Synthetic Minority Over-sampling Technique. *Journal of Artificial Intelligence Research* 16 (2002), 321–357. doi:10.1613/jair.953

[8] Tianqi Chen and Carlos Guestrin. 2016. XGBoost: A Scalable Tree Boosting System. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. 785–794. doi:10.1145/2939672.2939785

[9] Crunchbase. 2023. Crunchbase Data Methodology and Coverage. https://about.crunchbase.com/partners/data-research/ Accessed: 2026-02-15.

[10] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of NAACL-HLT*. 4171–4186. doi:10.18653/v1/N19-1423

[11] Failory. 2023. Startup Failure Rate: How Many Startups Fail and Why. https://www.failory.com/blog/startup-failure-rate Accessed: 2026-03-01.

[12] Tom Fawcett. 2006. An Introduction to ROC Analysis. *Pattern Recognition Letters* 27, 8 (2006), 861–874. doi:10.1016/j.patrec.2005.10.010

[13] Jerome H. Friedman. 2001. Greedy Function Approximation: A Gradient Boosting Machine. *Annals of Statistics* 29, 5 (2001), 1189–1232. doi:10.1214/aos/1013203451

[14] Paul Gompers and Josh Lerner. 2001. The Venture Capital Revolution. *Journal of Economic Perspectives* 15, 2 (2001), 145–168. doi:10.1257/jep.15.2.145

[15] Trevor Hastie, Robert Tibshirani, and Jerome Friedman. 2009. *The Elements of Statistical Learning* (2nd ed.). Springer. doi:10.1007/978-0-387-84858-7

[16] Haibo He and Edwardo A. Garcia. 2009. Learning from Imbalanced Data. *IEEE Transactions on Knowledge and Data Engineering* 21, 9 (2009), 1263–1284. doi:10.1109/TKDE.2008.239

[17] Steven N. Kaplan and Per Strömberg. 2009. Leveraged Buyouts and Private Equity. *Journal of Economic Perspectives* 23, 1 (2009), 121–146. doi:10.1257/jep.23.1.121

[18] William R. Kerr, Ramana Nanda, and Matthew Rhodes-Kropf. 2014. Entrepreneurship as Experimentation. *Journal of Economic Perspectives* 28, 3 (2014), 25–48. doi:10.1257/jep.28.3.25

[19] Ralph Kimball and Margy Ross. 2013. *The Data Warehouse Toolkit: The Definitive Guide to Dimensional Modeling* (3rd ed.). John Wiley & Sons.

[20] Arunesh Krishna, Ankur Agrawal, and Alok Choudhary. 2016. Predicting the Outcome of Startups: Less Failure, More Success. *Proceedings of IEEE 16th International Conference on Data Mining Workshops* (2016), 798–805. doi:10.1109/ICDMW.2016.0118

[21] Scott M. Lundberg and Su-In Lee. 2017. A Unified Approach to Interpreting Model Predictions. In *Advances in Neural Information Processing Systems*, Vol. 30. https://papers.nips.cc/paper/7062-a-unified-approach

[22] Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S. Corrado, and Jeff Dean. 2013. Distributed Representations of Words and Phrases and their Compositionality. In *Advances in Neural Information Processing Systems*, Vol. 26. https://papers.nips.cc/paper/5021-distributed-representations

[23] Ramana Nanda and Matthew Rhodes-Kropf. 2017. Financing Risk and Innovation. *Management Science* 63, 4 (2017), 901–918. doi:10.1287/mnsc.2015.2350

[24] Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. GloVe: Global Vectors for Word Representation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. 1532–1543. doi:10.3115/v1/D14-1162

[25] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. 3982–3992. doi:10.18653/v1/D19-1410

[26] Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. 2016. "Why Should I Trust You?": Explaining the Predictions of Any Classifier. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. 1135–1144. doi:10.1145/2939672.2939778

[27] Boris Sharchilev, Michael Roizner, Andrey Rumyantsev, and Denis Ozornin. 2018. Web-based Startup Success Prediction. *Proceedings of the 27th ACM International Conference on Information and Knowledge Management* (2018), 2283–2291. doi:10.1145/3269206.3272011

[28] Amit Singhal. 2001. Modern Information Retrieval: A Brief Overview. *IEEE Data Engineering Bulletin* 24, 4 (2001), 35–43. http://singhal.info/ieee2001.pdf

[29] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention Is All You Need. In *Advances in Neural Information Processing Systems*, Vol. 30. https://papers.nips.cc/paper/7181-attention-is-all-you-need

[30] Guang Xiang, Zhigang Zheng, Mianwei Wen, Jason Hong, Carolyn Rose, and Cong Liu. 2012. A Supervised Approach to Predict Company Acquisition with Factual and Topic Features Using Profiles and News Articles on TechCrunch. In *Proceedings of the International AAAI Conference on Web and Social Media*. https://ojs.aaai.org/index.php/ICWSM/article/view/14282

[31] Yi Yang, Mark Christopher Siy Uy, and Allen Huang. 2020. FinBERT: A Pretrained Language Model for Financial Communications. *arXiv preprint arXiv:2006.08097* (2020). https://arxiv.org/abs/2006.08097